

Cameras

Steve Rotenberg

CSE168: Rendering Algorithms

UCSD, Spring 2014

Radiance Equation

$$L_r(\omega_r) = \int_{\Omega} f_r(\omega_i, \omega_r) L_i(\omega_i) \cos \theta_i d\omega_i$$

$$L_r(\omega_r) = \int_{\varphi=0}^{2\pi} \int_{\theta=0}^{\frac{\pi}{2}} f_r(\omega_i, \omega_r) L_i(\omega_i) \cos \theta_i \sin \theta_i d\theta_i d\varphi_i$$

$$L_r(\omega_r) \approx \frac{2\pi}{N} \sum^N f_r(\omega_i, \omega_r) L_i(\omega_i) \cos \theta_i$$

Camera Focus

Camera Focus

- So far, we have been simulating *pinhole* cameras with perfect focus
- Often times, we want to simulate more realistic camera lenses that blur objects that are not in focus
- We say that real lenses have a limited *depth of field*, where the depth of field refers to the zone that is in focus
- Sometimes, this blur is referred to as *defocus blur*



Defocus Blur

- Defocus blur can be a bad thing if the subject of the image is out of focus
- However it can sometimes be a good thing
- If the subject is in focus and the background is blurred, this can have the effect of drawing the attention to the subject while removing distractions from the background
- It can have a nice artistic effect if handled properly



Lenses


- In photography, the term *lens* refers to the whole optical system in front of the film
- This is generally made from several *lens elements*, which are the individual pieces of glass in the lens, plus the iris and any structural components
- Most modern lenses have at least 4 elements, and complex zoom lenses can have more than 10



Focal Plane

- With a typical lens, there is a plane in front of the camera that is in perfect focus- this is called the *focal plane*
- Things get blurry as they get closer to the camera or further away from the focal plane

Focal Plane



Aperture

- A camera *aperture* is an opening through which light travels
- A small aperture will lead to a sharper image and a large aperture will lead to a blurrier image
- Typically, in a real camera, the aperture size can be changed with an adjustable *iris*



Rendering Camera Focus

- To add camera focus blur to a ray tracer, we need to model the camera lens
- We could model the entire complex lens with multiple lens elements, lens coatings, and an iris. This is actually a fairly common approach in high end movies, where computer generated objects need to be integrated into live-action scenes
- However, we will examine a much simpler method
- This method requires adding two more parameters to the Camera class: Aperture and FocalPlane
- Aperture refers to the diameter of the lens and FocalPlane is the distance in front of the camera

Rendering Camera Focus

- Our existing approach to generating camera rays uses a virtual image plane, which is 1.0 unit in front of the camera. We first generate a point on the image plane and then generate a ray from the camera origin through the point
- To modify this, all we need to do is scale the virtual image plane distance to the focal plane, and then generate a ray origin by choosing a random point on a circular disk the size of the camera aperture


Distant Focal Plane



Medium Focal Plane




Close Focal Plane



Small Aperture



Large Aperture




Bokeh

- In recent years, the Japanese word *bokeh* has been adopted by English speaking photographers and computer graphics practitioners to refer to the artistic effect of defocus blur
- More specifically, the term refers to how the lens renders out-of-focus points of light



Point Spread Function

- The *point spread function* (PSF) of an optical system describes how an individual point of light in the scene will map to the image



Motion Blur

Motion Blur

- *Motion blur* refers to the blurring we see on fast moving objects
- Motion blur is generally a good thing, and can improve the perceived realism in animations
- Motion blur is sometimes called *temporal antialiasing* (at least that's what it's called in the computer graphics world), and it reduces the aliasing phenomenon known as *strobing*

Shutter Speeds

- Motion blur occurs because the camera shutter is open for a finite length of time
- Camera shutter speeds vary based on light levels, exposure settings, film speed, etc.
- Typical shutter speeds range from $1/30^{\text{th}}$ of a second down to $1/4000^{\text{th}}$ of a second
- Some still images use very long shutter speeds (maybe a few seconds)
- Motion picture cameras (video or film) require the shutter speed to be faster than the frame time, so for 60Hz video, a typical shutter speed would be $1/100^{\text{th}}$ of a second or less
- Older film cameras typically run at 24 fps, and the shutter is typically open for up to half of the frame time, so $1/48^{\text{th}}$ of a second or less

Rendering Motion Blur

- To add motion blur to a ray tracer, we will need to distribute rays *in time*
- We add a 'Time' field to the Ray class
- When the camera generates a ray, we assign a random time
- We can either base it on actual time in seconds, or we can normalize it to a $[0...1]$ range
- When we trace the ray, we need to intersect it with objects moved to their correct position, based on the ray time

Moving Objects

- Not all objects in the scene need to move, so we can treat moving objects as a special case
- Just like we used the InstanceObject to position an object with a matrix, we can create a MotionObject which handles moving objects
- The MotionObject is a lot like an InstanceObject except it allows the matrix to change over time
- A simple way to do this is give it an initial and final matrix
- In the MotionObject::Intersect() function, we first use the input ray time to interpolate between the initial and final matrix. Then we have to compute the inverse on the fly, and from there, it behaves like a normal InstanceObject
- A more complex implementation could do an animation channel lookup and compute a matrix based on that

Matrix Interpolation

- Assuming we go with the simpler option, we still have to address the issue of how we interpolate between the initial and final matrix
- The simplest way is to just do a linear interpolation (lerp) for each component of the matrix

$$\text{Lerp}(t,a,b) = (1-t)a + tb = a + t(b-a)$$

- This will work reasonably well, assuming the object doesn't rotate too much in the time interval, which is usually the case

Matrix Interpolation

- However, for fast rotating objects (like a propeller), this may not be good enough
- When the matrix is linearly interpolated, every part of the object will move in a straight line from the initial to final position
- This is OK if the object only rotates a few degrees, but will start to break down if there is more than say 20 or so degrees of movement
- To improve on this, we can use quaternion interpolation or twist extraction

Matrix Twist Extraction

- This function extracts a rotation axis and angle from an orthonormal matrix

```
float Matrix34::ExtractTwist(Vector3 &out) const {  
    float theta=0.5f*sqrtf(1.0f + a.x*a.x + b.y*b.y + c.z*c.z);  
    if(theta<0.0000005f || theta>0.9999995f) {  
        if(theta<0.5f) printf("ERROR: Can't extract twist vector\n");  
        out.Set(1.0f,0.0f,0.0f);  
        return 0.0f;  
    }  
    float tmp=0.25f/theta;  
    out.Set(tmp*(b.z-c.y),tmp*(c.x-a.z),tmp*(a.y-b.x));  
    out.Normalize();  
    return 2.0f*acosf(theta);  
}
```

Matrix Interpolation

- To set up: (only need to do this once)
Matrix34 delta=InitialMatrix;
delta.Inverse();
delta.Dot(delta,FinalMatrix);
Angle=delta.ExtractTwist(Axis);
- To interpolate:
Matrix34 mtx;
mtx.MakeRotateUnitAxis(Axis,ray.Time*Angle);
mtx.Dot(InitialMatrix,mtx);
mtx.d.Lerp(ray.Time,InitialMatrix.d,FinalMatrix.d);
- Note: this process assumes orthonormal matrices


Non-Orthonormal Matrices

- For non-orthonormal matrices, it might be best to just lerp the components
- Alternatively, one can extract twist, shear, and scale, and then interpolate all of those and reconstruct a matrix, but this might be overkill

MotionObject

- Handling moving objects using this scheme is quite simple, and doesn't have a large performance penalty, apart from the necessity to shoot enough rays to reduce the noise in the blurred areas
- One catch to remember though, is that if you allow Objects to be placed in spatial data structures, you will need to compute a bounding box for the MotionObject that encompasses the object for the entire time interval


Moving Objects



Moving Cameras

- Objects are not the only thing that moves
- Often, the camera moves as well
- This may result in blurring everything during fast camera moves or turns
- The camera can be handled much like an object
- It can have an initial and final matrix and that can get interpolated as well
- The camera ray time is chosen first, then the camera matrix is interpolated, and then the ray origin and direction are built from the interpolated camera matrix

Moving Camera



Fast Panning Shot



Reflections & Shadows

- Motion blur comes from averaging across a finite interval of time
- Each camera ray is meant to be a single instant of that time interval
- Therefore, any reflected rays or shadow rays spawned from the initial camera ray **must use the same time as the camera ray**

Animation Blur

- Technically, anything that changes over time can be blurred
- This isn't limited just to matrices
- For example, camera FOV angles can change - a fast zoom-in will be blurred
- One could even blur changes in lighting properties (position, color, brightness...) or any other dynamic property



Lens Imperfections

Camera Imperfections

- Real camera lenses aren't perfect and can suffer from imperfections or *aberrations*
- Like defocus and motion blur, these can sometimes be desirable and sometimes undesirable
- However, if our goal is to model a real lens or integrate synthetic objects into a real scene, we may want to include some of these


Spherical Aberration

- Spherical aberration occurs when the individual rays coming from a point in the scene do not converge to a point on the film plane




Coma Aberration

- *Coma aberration* refers to the lens distortion that can cause off-axis point sources to appear to have a tail like a comet




Chromatic Aberration

- Chromatic Aberration is caused by the fact that the index of refraction for a lens varies with the wavelength of the light
- It can cause *color fringing* especially towards the edges of the image




Astigmatism

- Astigmatism is caused by lens elements that are not radially symmetric
- Rays in different planes focus to different points, leading to asymmetric distortions in the final image




Radial Distortion

- Lenses can cause various types of geometric distortion of the image
- Fish-eye lenses take advantage of this effect



Barrel distortion



Pincushion distortion



Fisheye lens (barrel)

Vignette

- *Vignetting* is the reduction in brightness or color saturation towards the edge of the image



Lens Flares

- Lens flares are caused by interreflection and scattering between the different elements and other components in a lens
- Very bright light sources tend to cause flares even if they are outside the image frame



Bloom, Halos, & Stars

- Bloom, halos, and stars are forms of lens flares



Modeling Lens Imperfections

- Most of the lens imperfections shown can be faked as a 2D post process
- We will discuss this in some more detail when we talk about HDR (high dynamic range) imaging in a later lecture
- There are some modern techniques however, that attempt to fully model the lens and capture all of these effects purely from the shape and arrangement of the lens elements

Camera Research Papers

- “A realistic camera model for computer graphics”, Kolb, Mitchell, Hanrahan, 1995
- “Polynomial optics: a construction kit for efficient ray-tracing of lens systems”, Hullin, Hanika, Heidrich, 2012
- “Efficient Monte Carlo rendering with realistic lenses”, Hanika, Dachsbacher, 2014

