

# Adaptive Sampling

Steve Rotenberg

CSE168: Rendering Algorithms

UCSD, Spring 2014


# Monte Carlo Integration

- *Monte Carlo Integration* is the process of using random numbers to estimate the value of an integral

# Estimating $\pi$

- Let's say that we want to use random Monte-Carlo integration to compute the value of pi
- We can do so by generating random 2D points in the  $[-1...1]$  interval and then counting how many of them are within a distance of 1.0 from the origin
- The ratio of the number points inside the circle to the total number of points will be approximately equal the ratio of the area of the circle ( $\pi$ ) to the area of the square (4)

$$\frac{N_{inside}}{N_{total}} \approx \frac{\pi}{4} = 0.78539816 \dots$$



# Adaptive Sampling

# Mean

- We can think of the previous process as estimating the value of a function by taking the average value of many estimates (or *samples*) of the function
- In the previous case, the value of each estimate was either 1.0 (inside) or 0.0 (outside), and they averaged out to  $\pi/4$
- Given a set of  $N$  estimates  $x_i$ , we can compute the mean  $\bar{x}$ :

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$


# Variance

- We can also compute the *variance* of a set of samples
- The variance is the average of the squared sample deviations from the mean
- The variance measures how far a set of numbers is spread out
- A variance of 0 indicates that all the samples are identical. A low value indicates they will tend to be spread over a small area near the mean, and a large value indicates they will be spread over a large range

$$\text{var}(x) = \sigma^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2$$

# Standard Deviation

- Standard deviation  $\sigma$  is the square root of variance
- Its advantage is that it is in the same units as the original estimates, so if the  $x_i$  values are in percentages, for example, then  $\sigma$  is a percentage
- If the original set of samples has a Gaussian distribution, then we would expect that  $\sim 68.2\%$  of the samples would be within  $\pm \sigma$  of the mean
- Keep in mind, that our  $\pi$  estimation example wasn't a Gaussian distribution, and most of our ray tracing applications won't be either, but that's OK



# Sample Standard Deviation

- The previous variance estimate is based on the idea that we have *all* of the relevant members of a set of numbers
- However, we're usually just working with a subset of possible values out of a much larger (or infinitely large) set
- To correct for this, it is better to use the *sample standard deviation*:

$$s = \sqrt{\frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2}$$



# Standard Error

- The *standard error of the mean* or just *standard error* is a probabilistic bound on the estimation of the mean

$$err = \frac{s}{\sqrt{N}}$$

- If we take 100 samples of some function and compute the mean, we want to know how accurate our estimate is of the ‘true mean’
- The standard error tells us the expected error range, but it isn’t a guarantee
- Like the standard deviation, it is based on the assumption of a Gaussian distribution of samples
- Plus, it is not an absolute bound but a probabilistic one
- Still, it is a reasonable way to compute the expected error on many practical sample distributions

# Error Estimation

- So now, we have a formula to estimate the error in our calculation of the mean

$$err = \sqrt{\frac{1}{N^2 - N} \sum_{i=1}^N (x_i - \bar{x})^2}$$

# Adaptive Sampling

- The idea behind *adaptive sampling* is to use the error estimate to determine when we've taken enough samples to reduce the error below some desired threshold
- To apply this to rendering, we can apply this on a pixel by pixel basis
- For each pixel, we start shooting rays, keeping track of the mean and variance of our estimate as we go, and using this to estimate the error
- When the error drops below some specified tolerance, we are done with that pixel
- It is best to start with a fixed number of samples (at least 10, but probably more) so that our initial error estimate is reasonable

# Running Counts

- From a look at the variance formula, we see that we need to know the mean in order to compute the variance
- This would imply that we need to store all of the values of  $x_i$  to compute the mean and variance at any point
- However, by rearranging the equation, we can put it into a form that allows us to just keep running counts of the mean of  $x_i$  and mean of  $x_i^2$  and adjust those with each new sample

$$\frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2 = \frac{1}{N} \sum_{i=1}^N x_i^2 - \left( \frac{1}{N} \sum_{i=1}^N x_i \right)^2$$

# Running Counts

- This way, we don't need to store all the previous samples, and we just need to track two numbers (let's call them  $m$  (for mean) and  $s$  (for square))
- For each sample  $i$ , we compute the value of the function  $f_i$  and then update  $m$  and  $s$

$$\begin{aligned} m &+= f_i \\ s &+= f_i^2 \end{aligned}$$

- At any time, we can then compute the mean variance as:

$$\begin{aligned} \text{mean} &= m/N \\ \text{var} &= s/N - \text{mean}^2 \end{aligned}$$

# Running Counts

- The sample standard deviation is

$$s = \sqrt{\frac{N}{N-1} \text{var}}$$

- And the error is

$$err = \sqrt{\frac{1}{N-1} \text{var}}$$

# Vectors

- As we are making estimates on colors, we really want to use vectors in this process
- In this case, our running count  $m$  is actually just a vector (color) and our running count for  $s$  is a scalar that sums up the magnitude squared
- The variance requires the mean squared, so we use the magnitude squared of the vector mean
- The variance, standard deviation, and error will therefore just be scalars

# Adaptive Sampling





# Adaptive Sampling Density



# Adaptive Sampling

- Adaptive sampling is a nice technique that will usually improve render times
- However, it can suffer from some problems
- If the initial set of samples is small, this has a tendency to underestimate the error, leading to premature termination of the sampling
- This is mainly due to the fact that rendering does not generate a Gaussian distribution of values

# Close-Up View



# Importance Sampling

# Mean Estimate

- We saw that when we are trying to estimate a function using Monte Carlo integration, our estimate after  $N$  samples is computed as:

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

# Weighted Average

- A slightly more sophisticated approach is to take a *weighted average*

$$\bar{x} = \frac{\sum w_i x_i}{\sum w_i}$$

- Is there some set of weights  $w_i$  that will give us a better estimate than just taking the mean?

# Importance Sampling

- The idea behind *importance sampling* is to attempt to select a weighted distribution of random numbers that will lead to a more accurate estimate of the function
- In order to do this however, we need to have some sort of idea of what the function might look like so we can select our weights

# Monte Carlo Integration

- Let's say that we are trying to estimate an integral over some domain
- We'll choose a 1D function  $f(t)$  over the domain  $[0...1]$  as an example, but this can be extended to n-D functions over any domain

$$z = \int_{t=0}^1 f(t) dt$$



# Monte Carlo Integration

$$z = \int_{t=0}^1 f(t) dt$$

- We want to estimate  $z$ , so we choose a set of  $N$  random numbers  $\xi_i$  on the  $[0...1]$  interval and evaluate  $f_i=f(\xi_i)$  for each one
- The estimate for the integral is then:

$$z \approx \frac{1}{N} \sum_{i=1}^N f(\xi_i)$$

# Monte Carlo Integration

- Without knowing anything at all about  $f()$ , we can't really do much better than that
- However, if we have some sort of idea of the shape of  $f()$ , then we may be able to choose a set of weighted samples that do a better job
- For example, let's say that we know that  $f()$  is actually the product of two other functions  $g()$  and  $h()$ :

$$f(t)=g(t)h(t)$$

- Let's also say that we know  $g(t)$  exactly, but have no idea what  $h(t)$  is

# Monte Carlo Integration

$$z = \int_{t=0}^1 f(t) dt = \int_{t=0}^1 g(t)h(t) dt$$

$$z \approx \frac{1}{N} \sum_{i=1}^N f(\xi_i) = \frac{1}{N} \sum_{i=1}^N g(\xi_i)h(\xi_i)$$

# Importance Sampling

- If we choose a set of random numbers that match the distribution of  $g()$ , then we can estimate the function with:

$$Z \approx \frac{\sum w_i h(\xi_i)}{\sum w_i}$$

- This will usually lead to a much better estimate, and allow us to take fewer samples
- This process is known as *importance sampling*, since we place more samples where we think they are more important (i.e., where we expect the function to have a higher value)
- The question that remains is how do we generate a set of random numbers that matches the distribution of  $g()$ ?

# Rejection Sampling

- We want to generate a bunch of samples that match the distribution of  $g(t)$  over the  $[0...1]$  interval
- One way to do this is with *rejection sampling*
- Let's say that the maximum value of  $g(t)$  over the interval is  $g_{max}$
- We start with a uniformly distributed random number  $\xi_i$ , but then choose a second random number  $\delta_i$
- If  $g(\xi_i)$  is less than  $g_{max} \delta_i$ , then we keep  $\xi_i$ , otherwise, we reject it and try again
- We keep this up until we have a  $\xi_i$  that passes
- This is then used to evaluate  $h(\xi_i)$  and contributes to our estimate
- This process may lead to a lot of rejections, which ultimately waste time, but this can be improved by using a tighter fitting bound than just  $g_{max}$

# Probability Density Function

- A more sophisticated way of generating these distributions uses the concept of a *probability density function* or *PDF*
- We start by normalizing  $g(t)$  so that it represents 100% of the distribution

$$G(t) = \frac{g(t)}{\bar{g}}, \quad \text{where} \quad \bar{g} = \int_{-\infty}^{\infty} g(t) dt$$

- $G(t)$  will match the shape of  $g(t)$ , but will be scaled so that:

$$\int_{-\infty}^{\infty} G(t) dt = 1$$

- Note that  $g(t)$  does not need to be smooth or continuous, but it does need to be positive for all values of  $t$

# Cumulative Distribution Function

- We can then compute the *cumulative distribution function* or *CDF*:

$$C(t) = \int_{-\infty}^t G(t)dt$$

- The graph of  $C(t)$  will go from 0 to 1 in a continuous and non-decreasing way as  $t$  increases
- Note that:

$$G(t) = \frac{d}{dt} C(t)$$

# Random Distributions

- To generate a random number with the desired distribution, we choose a uniformly distributed random number  $\xi_i$  and find out where  $C(\alpha_i) = \xi_i$
- $\alpha_i$  will then be a random number matching the original distribution of  $g()$
- We can then evaluate  $\bar{g}h(\alpha_i)$  and use it in our estimate



# Random Distributions

- This is what is happening when we perform the BRDF sampling, discussed in a previous lecture, and implemented in project 3
- For the Ashikhmin BRDF, this technique was used to compute the sampling direction
- By using this technique, we can create a bunch of samples that estimate our function, while allowing each sample to have equal weight
- Using this approach should significantly reduce the noise in situations where we have irregular distributions (such as the sharp highlights on the BRDF)